



# *Intel® Architecture and Tools*

## *JURECA - Tuning for the platform II*

**Presenter: Dr. Heinrich Bockhorst**

**Date: 18-05-2015**




# Agenda

- **Introduction**
- Processor Architecture Basics
- Composer XE
- Selected Intel® Tools

# Chips and Dies




- A chip is the package containing one or more dies (silicon)
- Major components of the die can be easily identified
- Example of a typical die:



“Die shot” of Intel® Core™ Processor

[Gordon Moore]



5/22/2015



# Parallelism

## Problem:

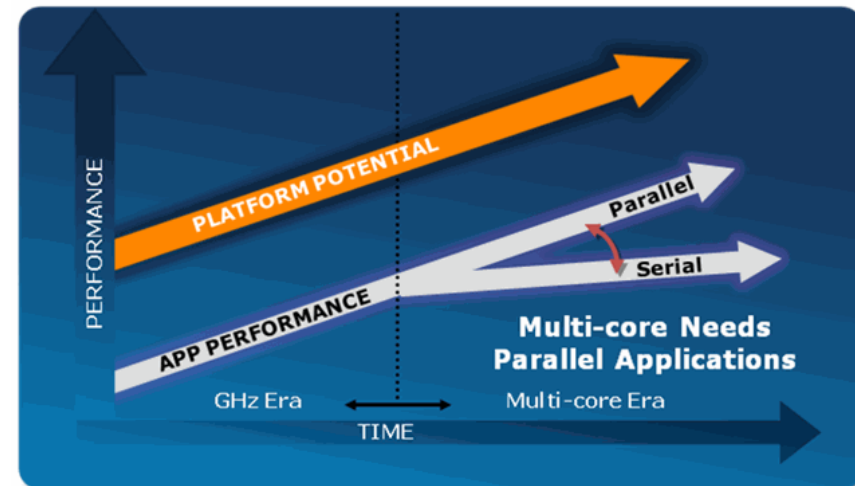
Economical operation **frequency of (CMOS) transistors is limited.**

⇒ **No free lunch anymore!**

## Solution:

More transistors allow more gates/logic on the same die space and power envelop, **improving parallelism:**

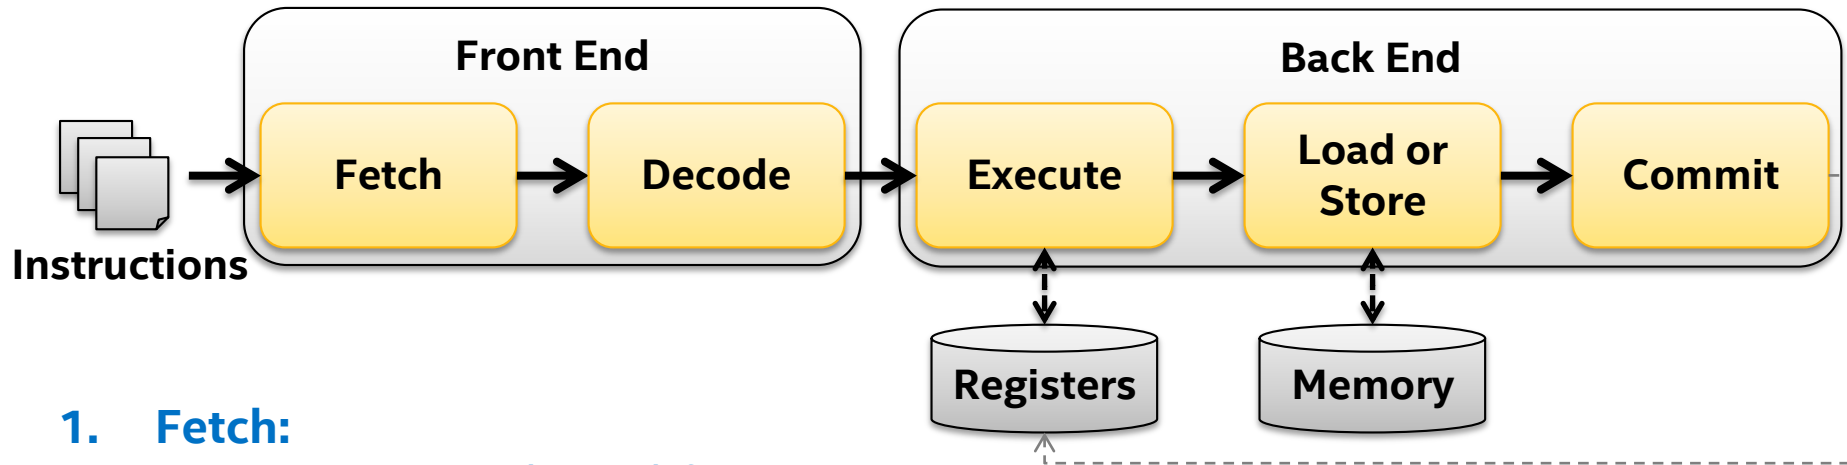
- **Thread level parallelism (TLP):**  
Multi- and many-core
- **Data level parallelism (DLP):**  
Wider vectors (SIMD)
- **Instruction level parallelism (ILP):**  
Microarchitecture improvements, e.g. threading, superscalarity, ...



# Processor Architecture Basics

## Pipeline

Computation of instructions requires several stages:



### 1. Fetch:

Read instruction (bytes) from memory

### 2. Decode:

Translate instruction (bytes) to microarchitecture

### 3. Execute:

Perform the operation with a functional unit

### 4. Memory:

Load (read) or store (write) data, if required


### 5. Commit:

Retire instruction and update micro-architectural state

5/22/2015

# Processor Architecture Basics

## Naïve Pipeline: Serial Execution



### Characteristics of strict serial execution:

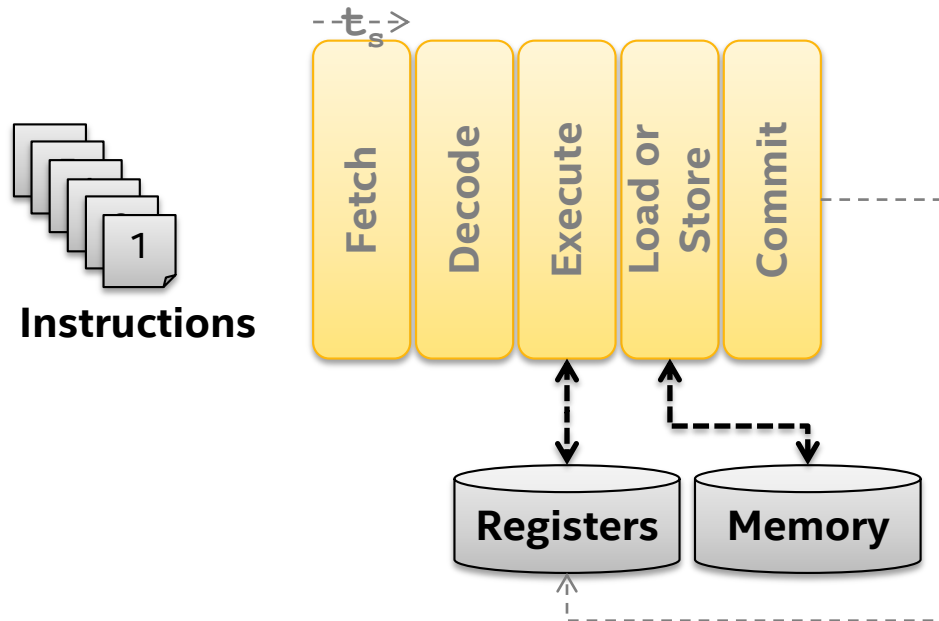
- Only one instruction for the entire pipeline (all stages)
- Low complexity
- Execution time:  $n_{\text{instructions}} * t_p$

### Problem:

- Inefficient because only one stage is active at a time

# Processor Architecture Basics

## Pipeline Execution



### Characteristics of pipeline execution:

- Multiple instructions for the entire pipeline (one per stage)
- Efficient because all stages kept active at every point in time
- Execution time:  $n_{\text{instructions}} * t_s$


### Problem:

- Reality check: What happens if  $t_s$  is not constant?



# Processor Architecture Basics

## Pipeline Stalls



### Pipeline stalls:

- Caused by pipeline stages to take longer than a cycle
- Caused by dependencies: order has to be maintained
- Execution time:  $n_{\text{instructions}} * t_{\text{avg}}$  with  $t_s \leq t_{\text{avg}} \leq t_p$

### Problem:

- Stalls slow down pipeline throughput and put stages idle.

# Processor Architecture Basics

## *Reducing Impact of Pipeline Stalls*

### **Impact of pipeline stalls can be reduced by:**

- Branch prediction
- Superscalarity + multiple issue fetch & decode
- Out of Order execution
- Cache
- Non-temporal stores
- Prefetching
- Line fill buffers
- Load/Store buffers
- Alignment
- Simultaneous Multithreading (SMT)

⇒ **Characteristics of the architecture that might require user action!**

# Processor Architecture Basics

## Superscalarity



### Characteristics of a superscalar architecture:

- Improves throughput by covering latency (ports are independent)
- Ports can have different functionalities (floating point, integer, addressing, ...)
- Requires multiple issue fetch & decode (here: 2 issue)
- Execution time:  $n_{\text{instructions}} * t_{\text{avg}} / n_{\text{ports}}$

### Problem:


- More complex and prone in case of dependencies

⇒ Solution: Out of Order Execution

5/22/2015

# Processor Architecture Basics

## Out of Order Execution



### Characteristics of out of order (OOO) execution:


- Instruction queue (I-Queue) moves stalling instructions out of pipeline
- Reorder buffer (ROB) maintains correct order of committing instructions
- Reduces pipeline stalls, but not entirely!
- Speculative execution possible
- Opposite of OOO execution is in order execution

# Processor Architecture Basics

## Alignment

### What happens if data spans across cache-lines?

- Loads/stores require all cache-lines of a datum to be loaded into cache.
- In worst case this doubles load/store times:



```
struct {  
    ... // 62 byte  
    int x; // 4 byte  
} data;
```

```
STRUCTURE /data/  
    ... ! 62 byte  
    INTEGER id ! 4 byte  
END STRUCTURE
```

Cache-line n

Cache-line n+1

Both cache-lines **n** and **n+1** need to be loaded up to L1 cache to access **x**!

Align data by padding or explicitly by compiler attributes to fit into exactly one cache-line, e.g.:

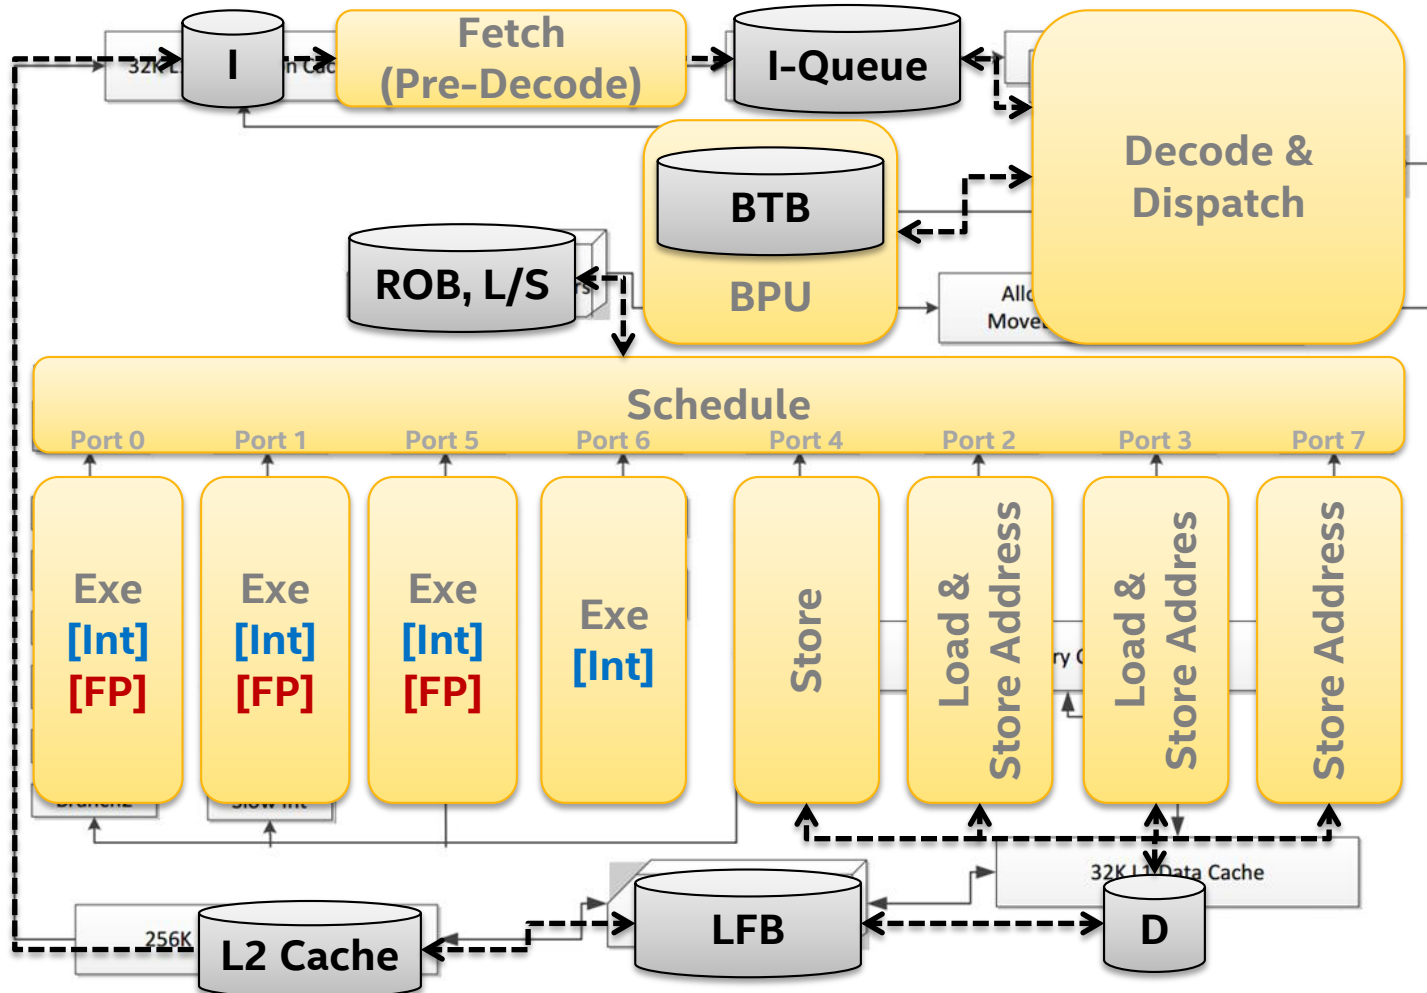
```
struct {  
    ... // 62 byte  
    ... // 2 byte  
    int x; // 4 byte  
} data;
```

```
STRUCTURE /data/  
    ... ! 62 byte  
    ... ! 2 byte  
    INTEGER id ! 4 byte  
END STRUCTURE
```

# Processor Architecture Basics

Example: 4<sup>th</sup> Generation Intel® Core™

From [Intel® 64 and IA-32 Architectures Optimization Reference Manual](#):



5/22/2015

# Processor Architecture Basics

## Core vs. Uncore

- **Core:**

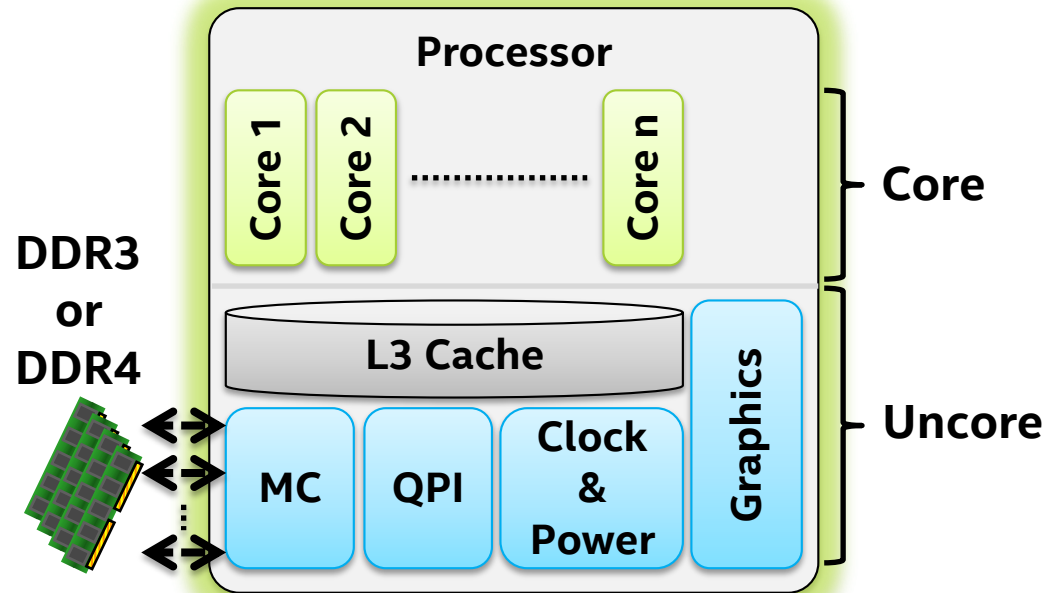
Processor core's logic:

- Execution units
- Core caches (L1/L2)
- Buffers & registers
- ...

- **Uncore:**

All outside a processor core:

- Memory controller/channels (MC) and Intel® QuickPath Interconnect (QPI)
- L3 cache shared by all cores
- Type of memory
- Power management and clocking
- Optionally: Integrated graphics

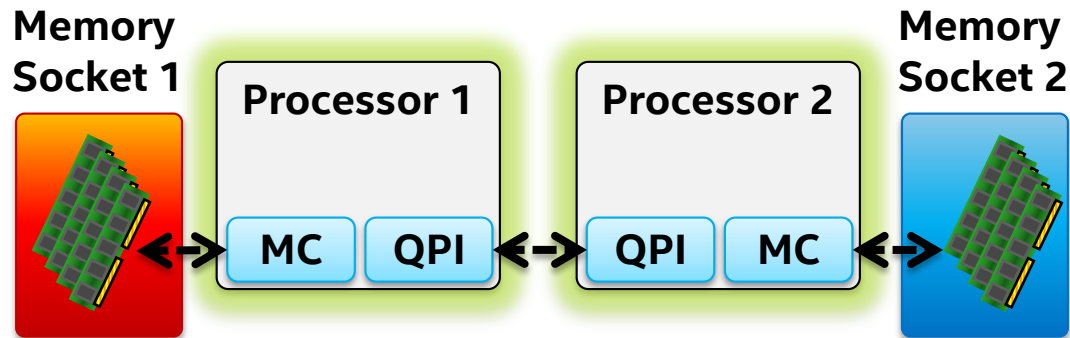


⇒ **Only uncore is differentiation within same processor family!**

5/22/2015

# Processor Architecture Basics

## UMA and NUMA



- **UMA (aka. non-NUMA):**

- Uniform Memory Access (UMA)
- Addresses interleaved across memory nodes **by cache line**
- Accesses may or may not have to cross QPI link
  - ⇒ **Provides good portable performance without tuning**



- **NUMA:**

- Non-Uniform Memory Access (NUMA)
- Addresses not interleaved across memory nodes by cache line
- Each processor has direct access to contiguous block of memory
  - ⇒ **Provides peak performance but requires special handling**





# Processor Architecture Basics

## NUMA - Thread Affinity & Enumeration

### Non-NUMA:

Thread affinity **might** be beneficial (e.g. cache locality) but not required

### NUMA:

Thread affinity is **required**:

- Improve accesses to local memory vs. remote memory
- Ensure 3<sup>rd</sup> party components support affinity mapping, e.g.:
  - Intel® OpenMP\* via `$KMP_AFFINITY` or `$OMP_PLACES`
  - Intel® MPI via `$I_MPI_PIN_DOMAIN` (default may be OK)  
tool `cpuinfo` provides additional information
  - ...
- Right way to get enumeration of cores:

### Intel® 64 Architecture Processor Topology Enumeration

<https://software.intel.com/en-us/articles/intel-64-architecture-processor-topology-enumeration>

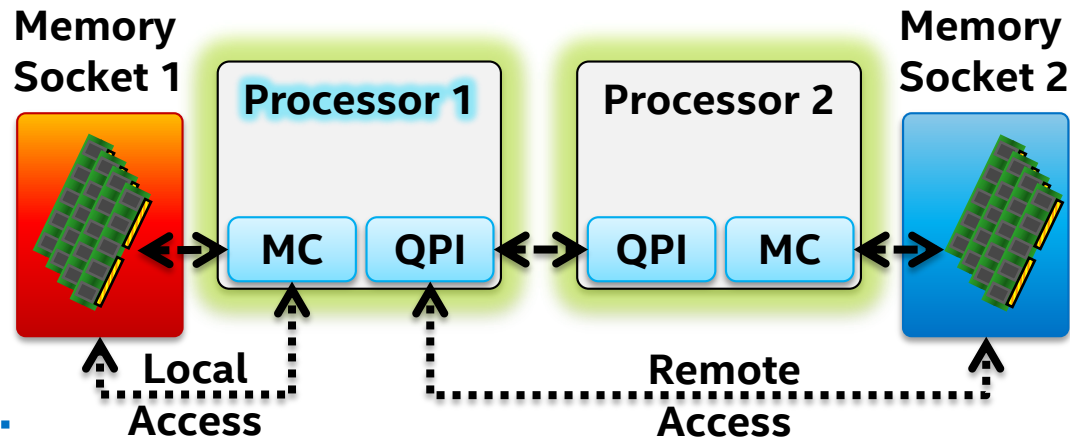
# Processor Architecture Basics

## NUMA - Memory, Bandwidth & Latency

### Memory allocation:

- Differentiate: implicit vs. explicit memory allocation
- Explicit allocation with NUMA aware libraries, e.g. **libnuma** (Linux\*)
- Bind **memory** ⇔ **(SW) thread**, and **(SW) thread** ⇔ **processor**
- More information on optimizing for performance:

<https://software.intel.com/de-de/articles/optimizing-applications-for-numa>



### Performance:

- Remote memory access **latency** **~1.7x** greater than local memory
- Local memory **bandwidth** can be up to **~2x** greater than remote

# Clock Rate and Power Gating

## Control power utilization and performance:

- Clock rate:
  - Idle components can be clocked lower:  
**Save energy**  
⇒ Intel SpeedStep®
  - If thermal specification allows, components can be over-clocked:  
**Higher performance**  
⇒ Intel® Turbo Boost Technology
  - Base frequency at P1; P0 might be turbo, depending on processor to decide (how many cores/GPU are active)
  - Intel® Turbo Boost Technology 2.0: processor can decide whether turbo mode can **exceed** the TDP with higher frequencies
- Power gating:  
Turn off components to save power
- P-state: Processor state; low latency; combined with speed step
- C-state: chip state; longer latency, more aggressive static power reduction

# Documentation

## **Intel® 64 and IA-32 Architectures Software Developer Manuals:**

- Intel® 64 and IA-32 Architectures Software Developer's Manuals
  - Volume 1: Basic Architecture
  - Volume 2: Instruction Set Reference
  - Volume 3: System Programming Guide
- Software Optimization Reference Manual
- Related Specifications, Application Notes, and White Papers

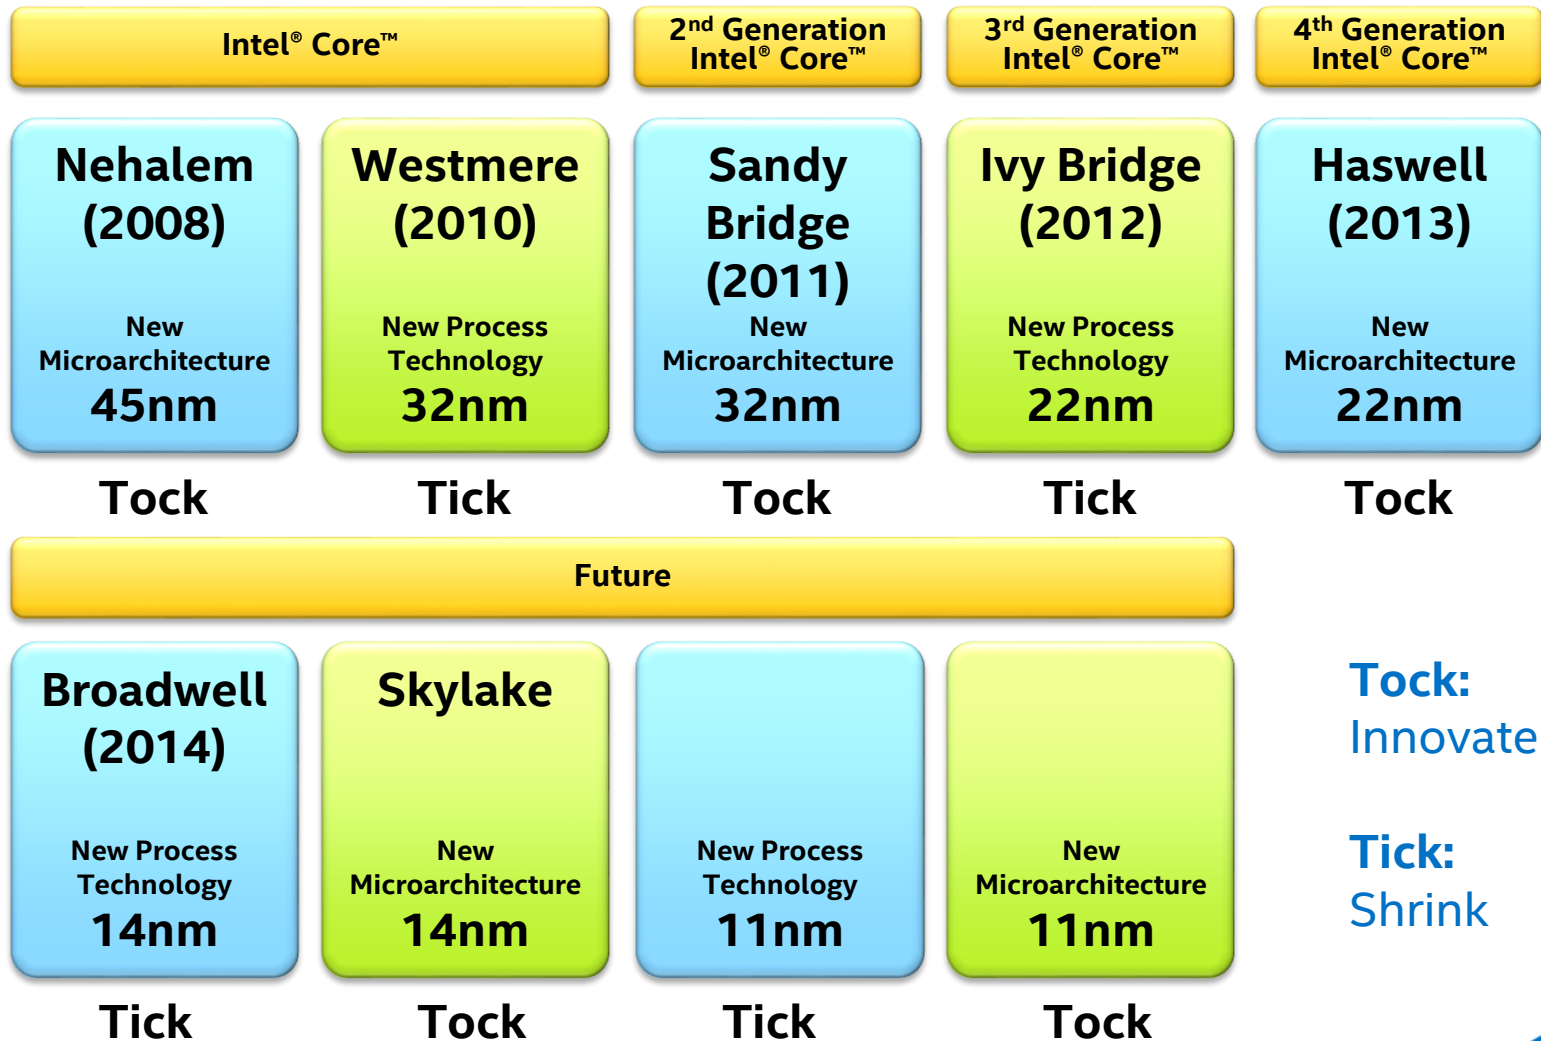

[https://www-ssl.intel.com/content/www/us/en/processors/architectures-software-developer-manuals.html?iid=tech\\_vt\\_tech+64-32\\_manuals](https://www-ssl.intel.com/content/www/us/en/processors/architectures-software-developer-manuals.html?iid=tech_vt_tech+64-32_manuals)

Intel® Processor Numbers (who type names are encoded):

[http://www.intel.com/products/processor\\_number/eng/](http://www.intel.com/products/processor_number/eng/)

# Desktop, Mobile & Server

## Tick/Tock Model



5/22/2015

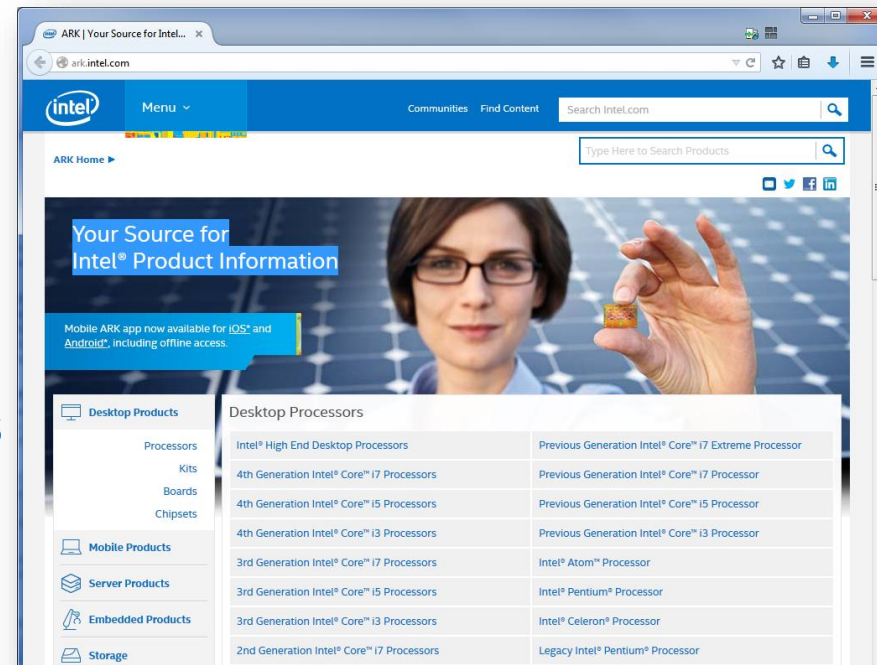
# Desktop, Mobile & Server

## Your Source for Intel® Product Information

### Naming schemes:

- Desktop & Mobile:
  - **Intel® Core™ i3/i5/i7** processor family
  - 4 generations, e.g.:  
4<sup>th</sup> Generation Intel® Core™ i7-XXXX
- Server:
  - **Intel® Xeon® E3/E5/E7** processor family
  - 3 generations, e.g.:  
Intel® Xeon® Processor E3-XXXX v3

Information about available Intel products can be found here: <http://ark.intel.com/>



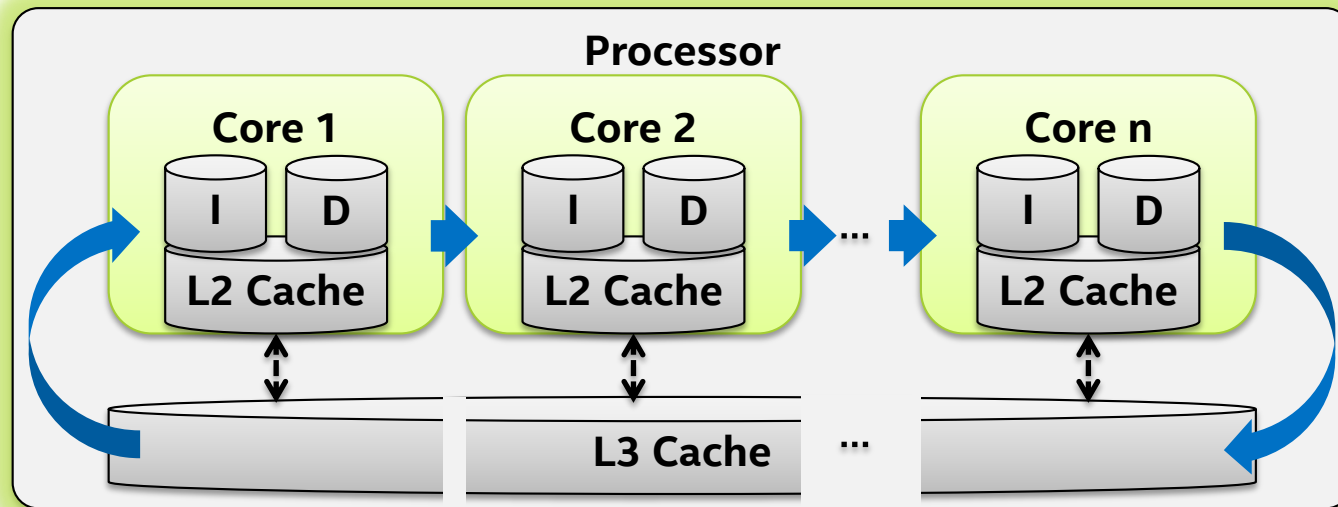
# Desktop, Mobile & Server

## Characteristics

- Processor core:
  - **4 issue**
  - Superscalar **out-of-order** execution
  - Simultaneous multithreading:  
Intel® Hyper-Threading Technology with **2 HW threads per core**
- Multi-core:
  - Intel® Core™ processor family: up to 8 cores (i7-5960) -- mobile desktop
  - Intel® Xeon® processor family: up to 18 cores (E5-2693 v3 ) – server  
all even numbers from 4-18 available
- Caches:
  - **Three level** cache hierarchy – L1/L2/L3 (Nehalem and later)
  - 64 byte cache line

# Desktop, Mobile & Server Caches

Cache hierarchy:



Level	Latency (cycles)	Bandwidth (per core per cycle)	Size
L1-D	4	2x 16 bytes	32KiB
L2 (unified)	12	1x 32 bytes	256KiB
L3 (LLC)	26-31	1x 32 bytes	varies ( $\geq$ 2MiB per core)
L2 and L1 D-Cache in other cores	43 (clean hit), 60 (dirty hit)		



# Desktop, Mobile & Server

## Intel® Hyper-Threading Technology II

- Not shared are:
  - Registers
  - Architectural state
- Smaller extensions needed:
  - More uops in backend need to be handled
  - ROB needs to be increased
- Easy and efficient to implement:
  - Low die cost: Logic duplication is minimal
  - Easy to handle for a programmer (multiple SW threads)
  - Can be selectively used, depending on workload
  - Many workloads can benefit from SMT – just enable it
- More insights to Intel® Hyper-Threading Technology:  
<https://software.intel.com/en-us/articles/performance-insights-to-intel-hyper-threading-technology>

5/22/2015

# Desktop, Mobile & Server Performance

- Following Moore's Law:

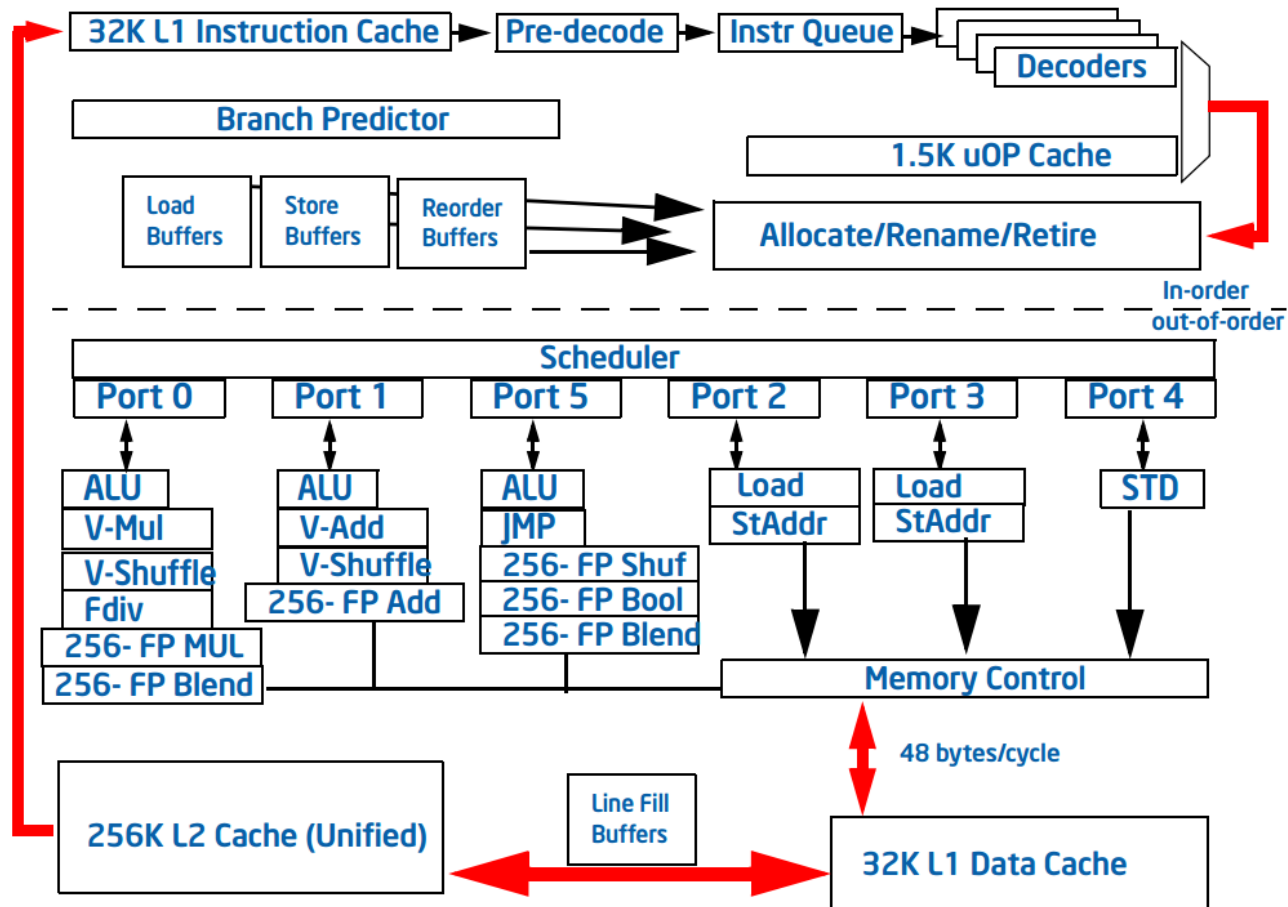
Microarchitecture	Instruction Set	SP FLOPs per Cycle per Core	DP FLOPs per Cycle per Core	L1 Cache Bandwidth (bytes/cycle)	L2 Cache Bandwidth (bytes/cycle)
Nehalem	SSE (128-bits)	8	4	32 (16B read + 16B write)	32
Sandy Bridge	Intel® AVX (256-bits)	16	8	48 (32B read + 16B write)	32
Haswell	Intel® AVX2 (256-bits)	32	16	96 (64B read + 32B write)	64

- Example of **theoretic peak FLOP** rates:
  - Intel® Core™ i7-2710QE (Sandy Bridge):  
 $2.1 \text{ GHz} * 16 \text{ SP FLOPs} * 4 \text{ cores} = \mathbf{134.4 \text{ SP GFLOPs}}$
  - Intel® Core™ i7-4765T (Haswell):  
 $2.0 \text{ GHz} * 32 \text{ SP FLOPs} * 4 \text{ cores} = \mathbf{256 \text{ SP GFLOPs}}$

# Desktop, Mobile & Server

## 2<sup>nd</sup> Generation Intel® Core™

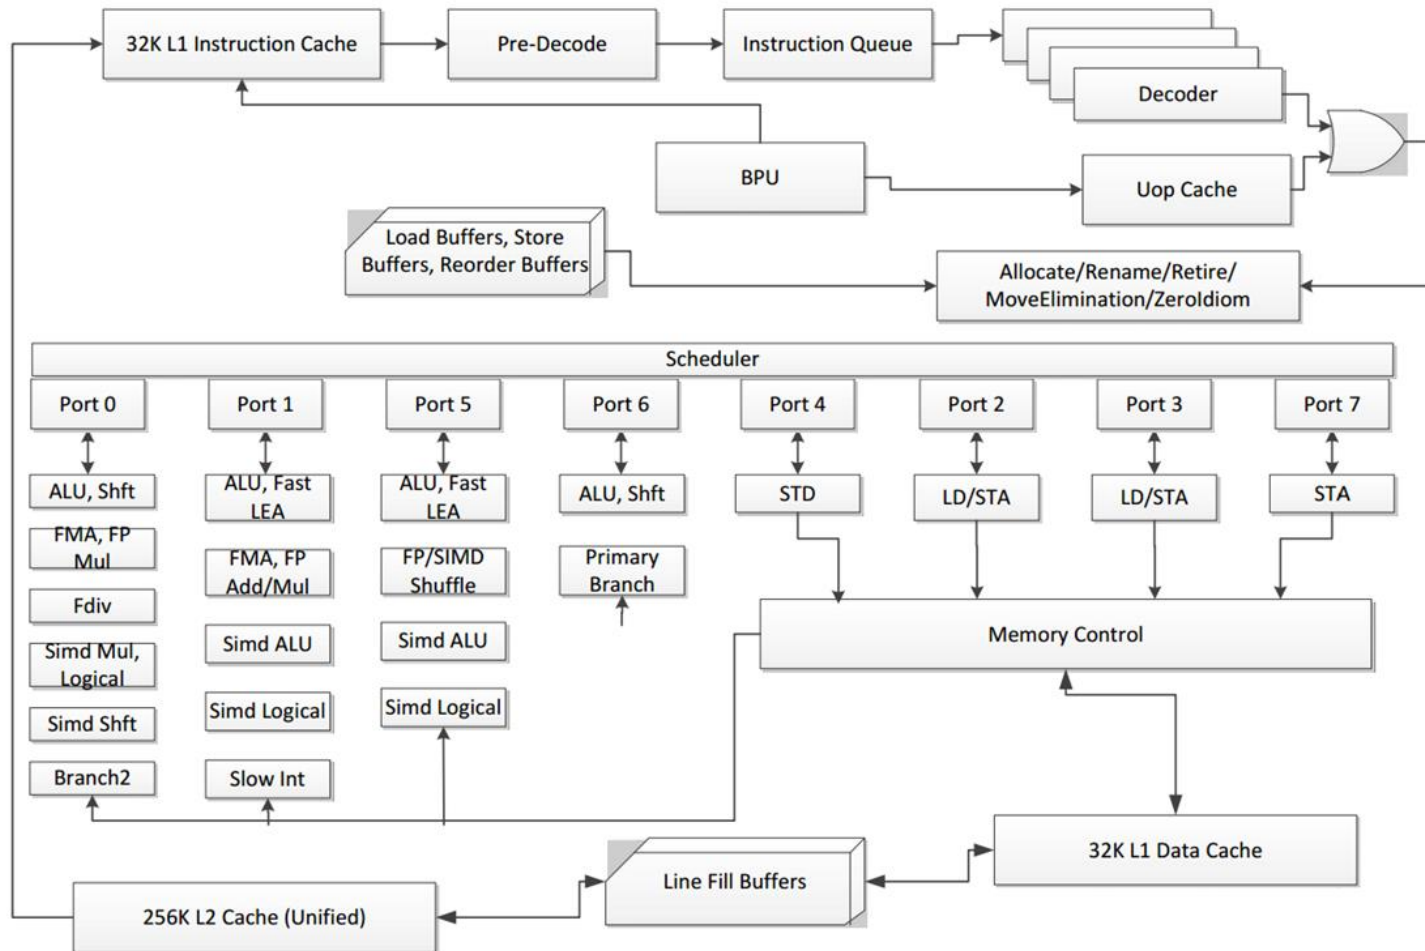
From [Intel® 64 and IA-32 Architectures Optimization Reference Manual](#):



# Desktop, Mobile & Server

## 4<sup>th</sup> Generation Intel® Core™

From [Intel® 64 and IA-32 Architectures Optimization Reference Manual](#):



5/22/2015



# Intel® Composer XE

# Common Optimization Options

	Windows*	Linux*, OS* X
Disable optimization	/Od	-O0
Optimize for speed (no code size increase)	/O1	-O1
Optimize for speed (default)	/O2	-O2
High-level loop optimization	/O3	-O3
Create symbols for debugging	/Zi	-g
Multi-file inter-procedural optimization	/Qipo	-ipo
Profile guided optimization (multi-step build)	/Qprof-gen /Qprof-use	-prof-gen -prof-use
Optimize for speed across the entire program ("prototype switch") <b>fast options definitions changes over time!</b>	/fast same as: /O3 /Qipo /Qprec-div-, /fp:fast=2 /QxHost)	-fast same as: <u>Linux</u> : -ipo -O3 -no-prec-div -static -fp- model fast=2 -xHost) <u>OS X</u> : -ipo -mdynamic-no-pic -O3 -no- prec-div -fp-model fast=2 -xHost
OpenMP support	/Qopenmp	-qopenmp
Automatic parallelization	/Qparallel	-parallel

# Vectorization

- **Single Instruction Multiple Data (SIMD):**
  - Processing vector with a single operation
  - Provides data level parallelism (DLP)
  - Because of DLP more efficient than scalar processing
- **Vector:**
  - Consists of more than one element
  - Elements are of same scalar data types (e.g. floats, integers, ...)
- **Vector length (VL):** Elements of the vector



# AVX Vector Types

## Intel® AVX



8x single precision FP



4x double precision FP

## Intel® AVX2



32x 8 bit integer



16x 16 bit integer



8x 32 bit integer



4x 64 bit integer



plain 256 bit



# Intel® AVX2 I

- Basically same as Intel® AVX with following additions:
  - **Doubles** width of **integer vector** instructions to 256 bits
  - Floating point fused multiply add (**FMA**)

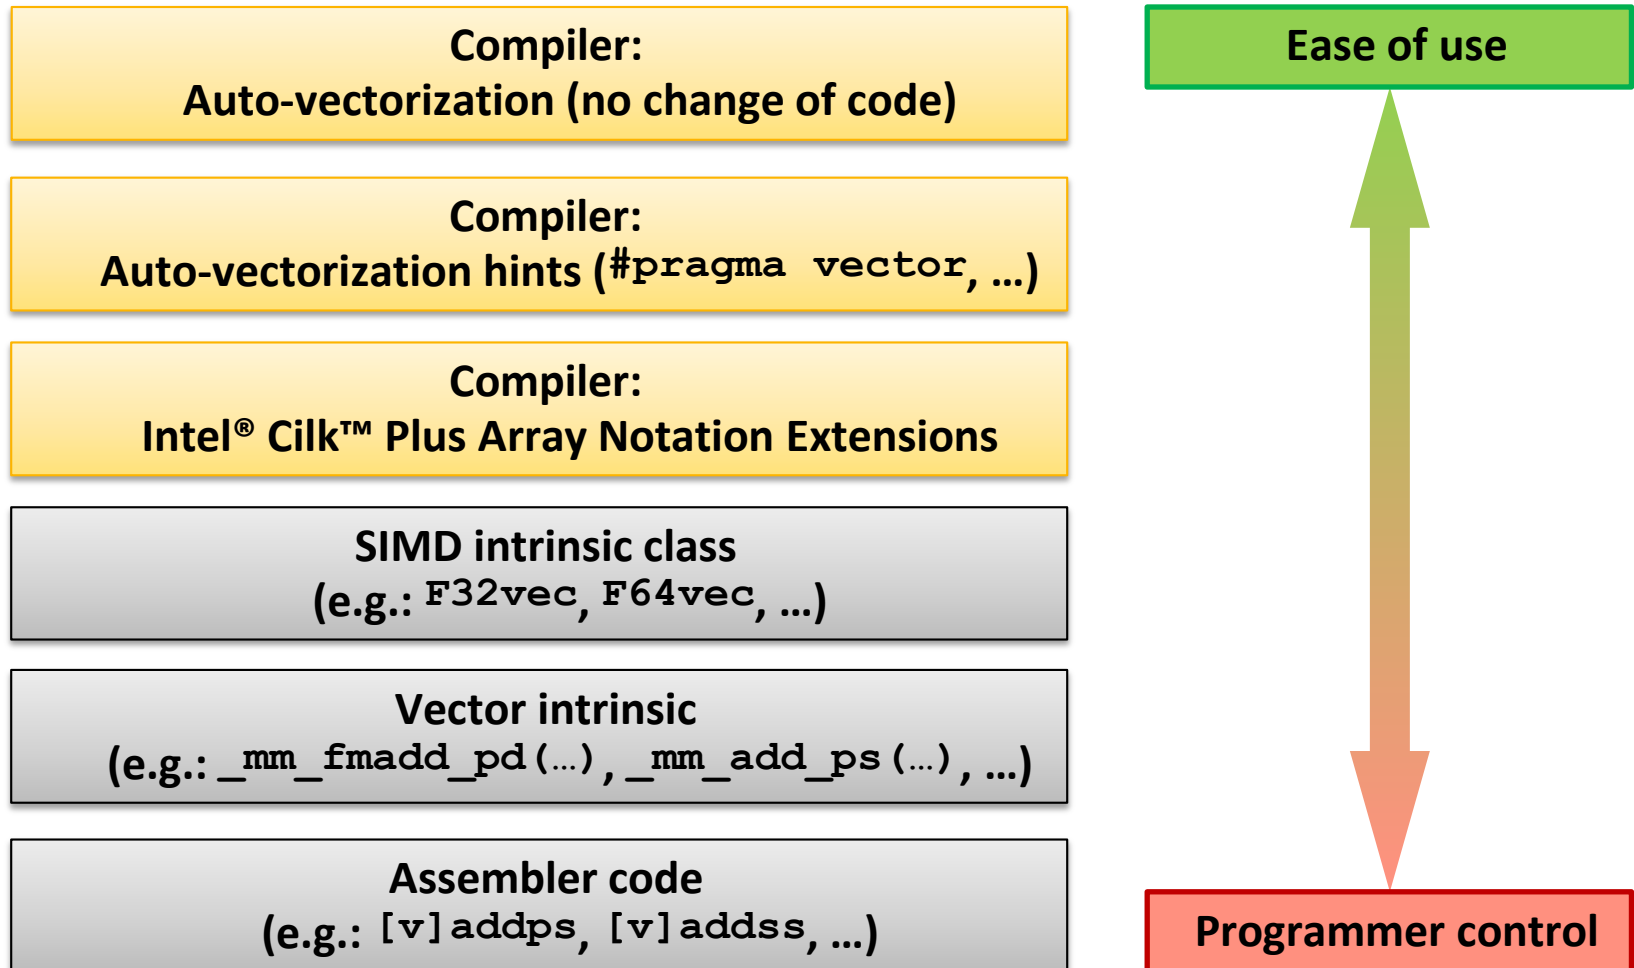
Processor Family	Instruction Set	Single Precision FLOPs Per Clock	Double Precision FLOPs Per Clock
Pre 2 <sup>nd</sup> generation Intel® Core™ Processors	SSE 4.2	8	4
2 <sup>nd</sup> and 3 <sup>rd</sup> generation Intel® Core™ Processors	AVX	16	8
<b>4<sup>th</sup> generation Intel® Core™ Processors</b>	<b>AVX2</b>	<b>32</b>	<b>16</b>

2x

4x

- Any-to-any permutes
- Vector-vector shifts

# Many Ways to Vectorize



# Basic Vectorization Switches I

- Linux\*, OS X\*: **-x<feature>**, Windows\*: **/Qx<feature>**
  - Might enable Intel processor specific optimizations
  - Processor-check added to “main” routine:  
Application errors in case SIMD feature missing or non-Intel processor with appropriate/informative message
- Linux\*, OS X\*: **-ax<features>**, Windows\*: **/Qax<features>**
  - Multiple code paths: baseline and optimized/processor-specific
  - Optimized code paths for Intel processors defined by **<features>**
  - Multiple SIMD features/paths possible, e.g.: **-axSSE2,AVX**
  - Baseline code path defaults to **-msse2 (/arch:sse2)**
  - The baseline code path can be modified by **-m<feature>** or **-x<feature>**  
(**/arch:<feature>** or **/Qx<feature>**)

# Basic Vectorization Switches II

- Linux\*, OS X\*: **-m<feature>**, Windows\*: **/arch:<feature>**
  - Neither check nor specific optimizations for Intel processors:  
Application optimized for both Intel and non-Intel processors for selected SIMD feature
  - Missing check can cause application to fail in case extension not available
- Default for Linux\*: **-msse2**, Windows\*: **/arch:sse2**:
  - Activated implicitly
  - Implies the need for a target processor with at least Intel® SSE2
- Default for OS X\*: **-xsse3** (IA-32), **-xssse3** (Intel® 64)
- For 32 bit compilation, **-mia32 (/arch:ia32)** can be used in case target processor does not support Intel® SSE2 (e.g. Intel® Pentium® 3 or older)

# Basic Vectorization Switches III

- Special switch for Linux\*, OS X\*: **-xHost**, Windows\*: **/QxHost**
  - Compiler checks SIMD features of current host processor (where built on) and makes use of latest SIMD feature available
  - Code only executes on processors with same SIMD feature or later as on build host
  - As for **-x<feature>** or **/Qx<feature>**, if “main” routine is built with **-xHost** or **/QxHost** the final executable only runs on Intel processors

# Vectorization Pragma/Directive

- SIMD features can also be set on a function/subroutine level via pragmas/directives:
  - C/C++:  
`#pragma intel optimization_parameter target_arch=<CPU>`
  - Fortran:  
`!DIR$ ATTRIBUTES OPTIMIZATION_PARAMETER:TARGET_ARCH= <CPU>`
- Examples:

- C/C++:

```
#pragma intel optimization_parameter target_arch=AVX
void optimized_for_AVX()
{
    ...
}
```

- Fortran:

```
function optimized_for_AVX()
!DIR$ ATTRIBUTES OPTIMIZATION_PARAMETER:TARGET_ARCH=AVX
    ...
end function
```

# Control Vectorization I

- Disable vectorization:
  - Globally via switch:  
Linux\*, OS X\*: **-no-vec**, Windows\*: **/Qvec-**
  - For a single loop:  
C/C++: **#pragma novector**, Fortran: **!DIR\$ NOVECTOR**
  - Compiler still can use some SIMD features
- Using vectorization:
  - Globally via switch (default for optimization level 2 and higher):  
Linux\*, OS X\*: **-vec**, Windows\*: **/Qvec**
  - Enforce for a single loop (override compiler efficiency heuristic) if semantically correct:  
C/C++: **#pragma vector always**, Fortran: **!DIR\$ VECTOR ALWAYS**
  - Influence efficiency heuristics threshold:  
Linux\*, OS X\*: **-vec-threshold[n]**  
Windows\*: **/Qvec-threshold[[:]n]**  
**n: 100** (default; only if profitable) ... **0** (always)

# Control Vectorization II

- Verify vectorization:
  - Globally:  
Linux\*, OS X\*: **-opt-report**, Windows\*: **/Qopt-report**
  - Abort compilation if loop cannot be vectorized:  
C/C++: **#pragma vector always assert**  
Fortran: **!DIR\$ VECTOR ALWAYS ASSERT**
- Advanced:
  - Ignore vector dependencies (IVDEP):  
C/C++: **#pragma ivdep**  
Fortran: **!DIR\$ IVDEP**
  - “Enforce” vectorization:  
C/C++: **#pragma simd**  
Fortran: **!DIR\$ SIMD**



# How to detect Vectorization?

- We need feedback! Compiling with a flag below reveals a short report showing which loops are vectorized and the reason for refusing vectorisation for certain loops:
  - **-opt-report=<n>**  
n = 0-5 increases the level of detail, 2 is the default
  - **-opt-report-phase=vec**  
shows just vectorization
- Profile the code, compile without vectorization (-no-vec) and profile again. Look for code sections with different timing
- Look into assembly. Can be done by using Intel® VTune™ Amplifier XE

# How to start?

- Compile with minimal options and check timing
- Compile with `-xHost` and `-opt-report=5` and check timing
- Compile with `-xHost` and `-no-vec` disables vectorization. Compare with previous timing
- Use loopprofiler for single threaded profile on loop basis:

`icc -profile-functions -profile-loops=all -profile-loops-report=2 ...`

view results with `loopprofilerviewer.sh <datafile>`

- New Book: **Optimizing HPC Applications with Intel Cluster Tools**

# *Other Intel® HPC tools*

Intel® MPI

Intel® VTune™ Amplifier XE

Intel® Trace Analyzer and Collector (ITAC)

Intel® Math Kernel Library (MKL)

Intel® MPI Benchmarks (IMB)

Intel® Inspector XE

Intel® Advisor

All Tools available in Intel® Cluster Studio XE (ICS)

<http://software.intel.com/en-us/articles/intel-cluster-studio-xe/>

# VTune™ Amplifier XE hotspots

**Source View / Per line localization**

The screenshot displays the Intel VTune Amplifier XE 2013 interface. The main window is titled "/home/michome/rreed/projects/matrix/linux/matrix-mic - Intel VTune Amplifier". The "Lightweight Hotspots" tab is active, showing a "Hotspots viewpoint (change)" with tabs for Analysis Target, Analysis Type, Summary, Bottom-up, Caller/Callee, and Top-down Tree. The "Source" tab is selected, showing C code for a matrix multiplication. A yellow box highlights the "Source View / Per line localization" text. A yellow arrow points from this text to the "Source" tab. The code is as follows:

```
113     for (k = k0; k < k0 + mblock; k++) {
114         #pragma unroll(8)
115         #pragma ivdep
116         for (j = j0; j < j0 + mblock; j++) {
117             c[i][j] = c[i][j] + a[i][k] * b[k][j];
118         }
119     }
120 }
121 }
122 }
123 }
124 #elif defined (USE_OMP) //{
125     #pragma omp parallel for collapse (2)
126     for(i=0; i<msize; i++) {
127         for(k=0; k<msize; k++) {
128             #pragma unroll(8)
129             #pragma ivdep
130             for(j=0; j<msize; j++) {
131                 c[i][j] = c[i][j] + a[i][k] * b[k][j];
132             }
133         }
134     }
135 #endif // }
136 }
137 }
```

The "Assembly" tab is also visible, showing assembly code for Block 23 and Block 24. The assembly code is as follows:

```
Block 23:
0x402a90 131 lea (%rsi,%r13,1), %r9 0.028s
0x402a94 131 vprefetch0 (%rcx) 0.183s
0x402a98 131 vbroadcastsq (%r9,%rbx,8), %k0, %zmm0 2.954s
0x402a9f 131 vprefetch0 0x200(%rcx) 0.174s
0x402aa6 131 vprefetch0 (%r8) 0.294s
0x402aab 130 movsxd %edx, %r9 0.018s
0x402aae 131 vprefetch0 0x200(%r8) 0.073s
0x402ab5 131 mov %al, %al 0.009s
0x402ab7 131 vprefetch0 0x40(%rcx) 0.248s
0x402abc 131 mov %al, %al 0.009s
0x402abe 131 vprefetch0 0x240(%rcx) 0.257s

Block 24:
0x402ac5 131 vmovapd (%r8,%rax,8), %k0, %zmm1 6.147s
0x402acc 131 vprefetch1 0x1000(%rcx,%rax,8) 11.817s
0x402ad4 131 vmovapd 0x40(%r8,%rax,8), %k0, %zmm2 6.982s
0x402adc 131 vprefetch0 0x400(%rcx,%rax,8) 9.000s
0x402ae4 131 vmovapd 0x80(%r8,%rax,8), %k0, %zmm3 6.596s
0x402aec 131 vprefetch1 0x1000(%r8,%rax,8) 13.835s
0x402af4 131 vmovapd 0xc0(%r8,%rax,8), %k0, %zmm4 6.917s
0x402afc 131 vprefetch0 0x400(%r8,%rax,8) 8.486s
```

The bottom of the window shows filters: "No filters are applied.", "Process: Any Process", "Thread: Any Thread", "Module: Any Module", "Utilization: Any Utilization", "Call Stack Mode: Only user functions", "Inline Mode: on", "Loop Mode: Functions only".

# *Intel® Inspector XE*

## *Threading Correctness*

- Inspector XE simulates all possible orders of data access
- Not only errors in this run. Inspector XE finds all possible errors even if they don't happen in this run!
- Typical race conditions in OpenMP will be due to forgotten variables in the private or reduction clause
- Following example shows a forgotten reduction clause
- Inspector can be used to find all necessary variables for the private/reduction clause!

# Inspector XE – Data Races

Intel Inspector XE 2013

Data race

Target Analysis Type Collection Log Summary Sources

Write - Thread OMP Worker Thread #2 (55255) (poisson-err.xlresiduum - compute.c:183)

compute.c

Source file not found. Code from cache is shown instead.

Call Stack

poisson-err.xlresiduum - compute.c:183

```
176
177 //pragma omp parallel for private(i,j) reduction(+:resid)
178 #pragma omp parallel for private(i,j)
179 for(i=1; i< n+1; i++)
180 for(j=1; j<m+1; j++)
181 {
182 double diff = x_new[i][j] - x_cur[i][j];
183 resid += diff*diff;
184 }
185
186 #ifdef USE_MPI
187 resid_tmp = resid;
188 MPI_Allreduce(&resid_tmp,&resid,1,MPI_DOUBLE,MPI_SUM,MPI_COMM_WORLD);
189 #endif
190
191 return resid;
```

Write - Thread main (55244) (poisson-err.xlresiduum - compute.c:183)

compute.c

Source file not found. Code from cache is shown instead.

Call Stack

poisson-err.xlresiduum - compute.c:183  
poisson-err.xlresiduum - compute.c:178  
poisson-err.xlmain - poisson.c:204  
poisson-err.xl\_start

```
176
177 //pragma omp parallel for private(i,j) reduction(+:resid)
178 #pragma omp parallel for private(i,j)
179 for(i=1; i< n+1; i++)
180 for(j=1; j<m+1; j++)
181 {
182 double diff = x_new[i][j] - x_cur[i][j];
183 resid += diff*diff;
184 }
185
186 #ifdef USE_MPI
187 resid_tmp = resid;
188 MPI_Allreduce(&resid_tmp,&resid,1,MPI_DOUBLE,MPI_SUM,MPI_COMM_WORLD);
189 #endif
190
191 return resid;
```

HINT: Synchronization allocation site - Thread main (55244) (poisson-err.xlcopy - compute.c:109)

Correct OpenMP  
pragma was  
commented

Both threads are  
writing and  
reading to/from  
variable resid

# Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS". NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

Copyright © 2014, Intel Corporation. All rights reserved. Intel, Pentium, Xeon, Xeon Phi, Core, VTune, Cilk, and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.

## Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

